Khaja Masroor Ahmed

Assignment 4

CS 751: Introduction to Digital Libraries
Dr. Michael Nelson
Spring 2015

May 2, 2015

Contents

1

Question 1

1.1 Question

- Using the pages from A3 that boilerpipe successfully processed, download those representations again & reprocess them with boilerpipe.
- Document the time difference (e.g., Time(A4) Time(A3)).
- Compute the Jaccard Distance x for each pair of pages (i.e., P(A3) & P(A4) for:
  - Unique terms (i.e., unigrams)
  - Bigrams
  - Trigrams
- See: `http://en.wikipedia.org/wiki/Jaccard_index`
- For each of the 3 cases (i.e., 1-, 2-, 3-grams) build a Cumulative Distribution Function that shows the % change on the x-axis & the % of the population on the x-axis
- See: `http://en.wikipedia.org/wiki/Cumulative_distribution_function`
- Give 3-4 examples illustrating the range of change that you have measured.

1.2 Solution

- The time difference between the time since I ran the boilerpipe for the URIs for assignment three and assignment four was 30 days.
- I used the same boilerpipe library as I used in assignment three.
- From the 10000 URIs I ran through boilerpipe I was able to successfully retrieve the data from 6086 URIs.
- I used the 'ngrams' python library for fetching the unigrams, bigrams and trigrams. I used the set data structure to ensure that only the unique terms were included in the comparison.
- The Jaccard Distance can be calculated by finding the union and the intersection. It is the ratio of the difference of the union and intersection by the union of the set.

- Below are the graphs for unigrams, bigrams and trigrams.

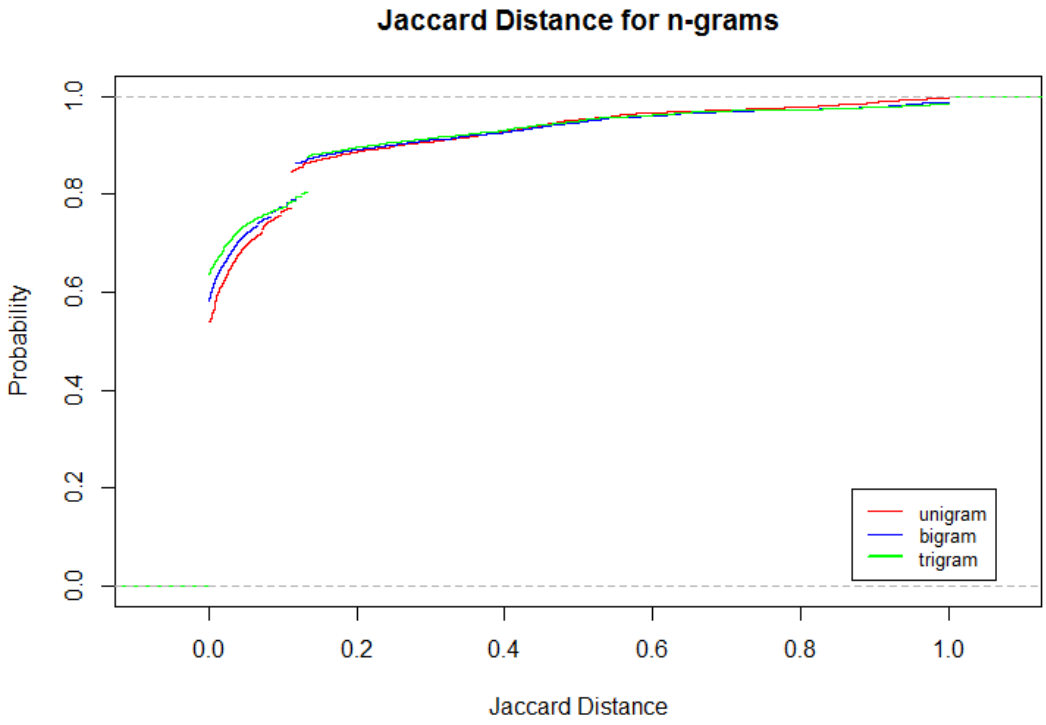

Fig. 1.1. CDF - Jaccard Distance vs. Probability

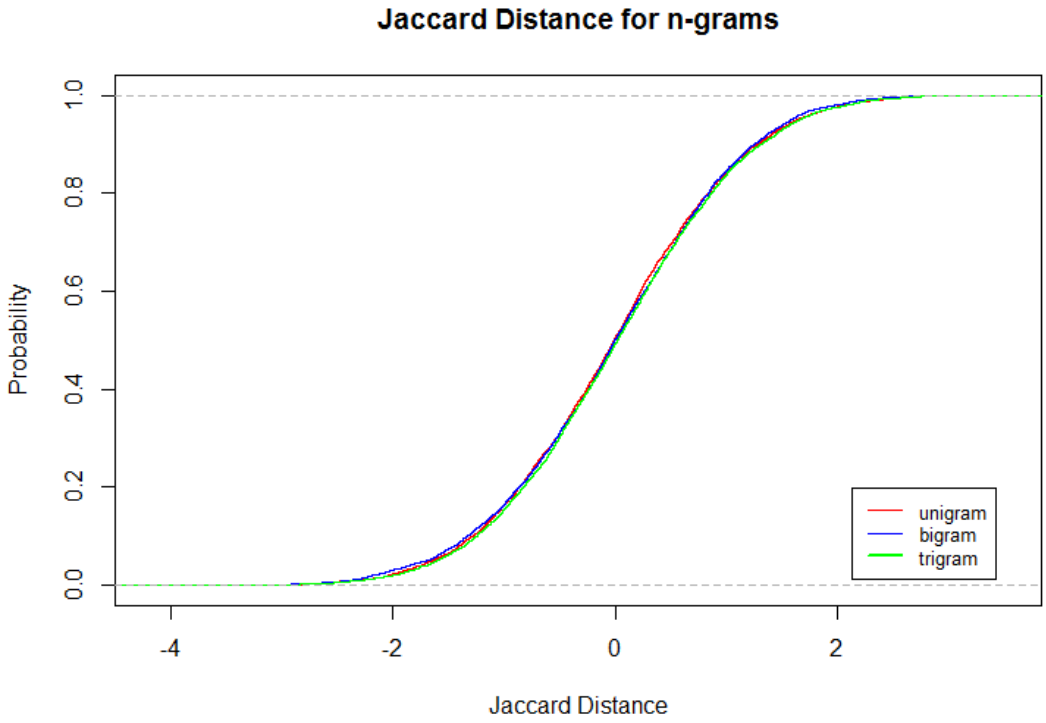

Fig. 1.2. CDF normalized - Jaccard Distance vs. Probability

1.3 Code Listing

```
1  import json
2  from boilerpipe.extract import Extractor
3  import datetime
4
5  f = open('status.txt', 'r+')
6  print datetime.datetime.now()
7  for line in f:
8          data = json.loads(line)
9          try:
10                 if data['index'] == 10001:
11                         print 'Program Executed'
12                         break
13                 finalURL=data['tweetURLData'][0]['finalURL']
14                 extractor = Extractor(extractor='
                       DefaultExtractor', url=finalURL)
15                 extracted_text = extractor.getText()
16                 link = str(data['index']) + '.txt'
17                 g = open(link, 'w')
18                 g.write(extracted_text.encode('utf-8'))
19                 g.close()
20         except:
21                 print datetime.datetime.now()
22                 print data['index']
23                 continue
```

Listing 1.1. Python program for fetching the boilerpipe content from the URIs.

```
from nltk.util import ngrams
import os

MAX_NGRAMS=3

def jaccardDistance(list1, list2):
        union = set(list1).union(list2)
        intersect = set(list1).intersection(list2)
        if len(union) == 0:
                return 0
        dist = (len(union) - len(intersect)) * 1.0 / len(
            union)
        return dist

def getNGrams(fileName):
        f = open(fileName, 'r+')
        unigramList=[]
        bigramList=[]
        trigramList=[]

        for line in f:
                unigram = ngrams(line.split(), MAX_NGRAMS-2)
                bigram = ngrams(line.split(), MAX_NGRAMS-1)
                trigram = ngrams(line.split(), MAX_NGRAMS)
                for grams in unigram:
                        unigramList.append(grams)
                #print len(unigramList)
                #print unigramList
                if len(unigramList) == 0:
                        print fileName
                        return unigramList,bigramList,
                            trigramList,True
                for grams in bigram:
                        bigramList.append(grams)
                for grams in trigram:
                        trigramList.append(grams)
        f.close()
        return unigramList, bigramList, trigramList, False

def fileExists(fileName):
        return os.path.isfile(fileName)
count = 1
uni = open('unigramOutput.txt','w')
bi = open('bigramOutput.txt','w')
tri = open('trigramOutput.txt','w')
while count < 10001:
        fileName1 = 'a3/' + str(count) + '.txt'
        fileName2 = 'a4/' + str(count) + '.txt'
```

```
47          if fileExists(fileName1) and fileExists(fileName2):
48                  unigramListA3=[]
49                  bigramListA3=[]
50                  trigramListA3=[]
51                  shouldExit = False
52                  unigramListA3, bigramListA3, trigramListA3,
                        shouldExit = getNGrams(fileName1)
53
54                  if shouldExit:
55                          print 'exit'
56                          count += 1
57                          continue
58
59                  unigramListA4=[]
60                  bigramListA4=[]
61                  trigramListA4=[]
62                  unigramListA4, bigramListA4, trigramListA4,
                        shouldExit = getNGrams(fileName2)
63
64                  if shouldExit:
65                          print 'exit2'
66                          count += 1
67                          continue
68                  #print str(jaccardDistance(unigramListA3,
                        unigramListA4))
69                  uni.write(str(jaccardDistance(unigramListA3,
                        unigramListA4)))
70                  uni.write('\n')
71                  bi.write(str(jaccardDistance(bigramListA3,
                        bigramListA4)))
72                  bi.write('\n')
73                  tri.write(str(jaccardDistance(trigramListA3,
                        trigramListA4)))
74                  tri.write('\n')
75          count += 1
76  uni.close()
77  bi.close()
78  tri.close()
79  print 'Program Executed'
```

Listing 1.2. Python program for calculating the Jaccard Distance.

```
dp1 <- read.table('c:/users/kahmed/desktop/unigramOutput.txt', header=FALSE)
dp2 <- read.table('c:/users/kahmed/desktop/bigramOutput.txt', header=FALSE)
dp3 <- read.table('c:/users/kahmed/desktop/trigramOutput.txt', header=FALSE)
datapoint1 <- dp1[,1]
datapoint2 <- dp2[,1]
datapoint3 <- dp3[,1]
X1 = rnorm(sort(datapoint1))
X2 = rnorm(datapoint2)
X3 = rnorm(datapoint3)
P1 = ecdf(datapoint1)
P2 = ecdf(X2)
P3 = ecdf(X3)
plot(P1, col="red", xlab="Number of Mementos", ylab="Probability", main = "Number of Mementos for each URI")
lines(P2, col="blue")
lines(P3, col="green")
legend("bottomright",inset = 0.05, c("unigram","bigram", "trigram"),
       cex=.8, col=c("red","blue", "green"), lwd=c(1,1.5,2))#pch=c(1,3))
```

Listing 1.3. R program for generating the CDF for Jaccard Distance.

2

Question 2

2.1 Question

- Using the pages from Q1 (A4), download all TimeMaps (including TimeMaps with 404 responses, i.e. empty or null TimeMaps)
  - Upload all the TimeMaps to github
- Build a CDF for # of mementos for each original URI (i.e., x-axis = # of mementos, y-axis = % of links)
- See: `http://timetravel.mementoweb.org/guide/api/`

2.2 Solution

- For downloading the TimeMaps I modified my script from the first assignment to retrieve the html pages. I just had to add a prefix to the URL to include the memento aggregator URL. But I wasn't able to successfully download the TimeMaps for most of the files.
- I then started using the script as provided in the mailing group by Alexander Nwala for retrieving the TimeMaps, but I noticed that the script was taking too long to run.
- Soon after I got back to using my script with a modified memento aggregator URL to point to the following URL `http://labs.mementoweb.org/timemap/json/`. I was successfully able to retrieve in JSON format which would make my life simpler in processing the data.
- Soon after I committed the TimeMaps to github.
- I wrote a script to find the number of mementos in each of the URIs

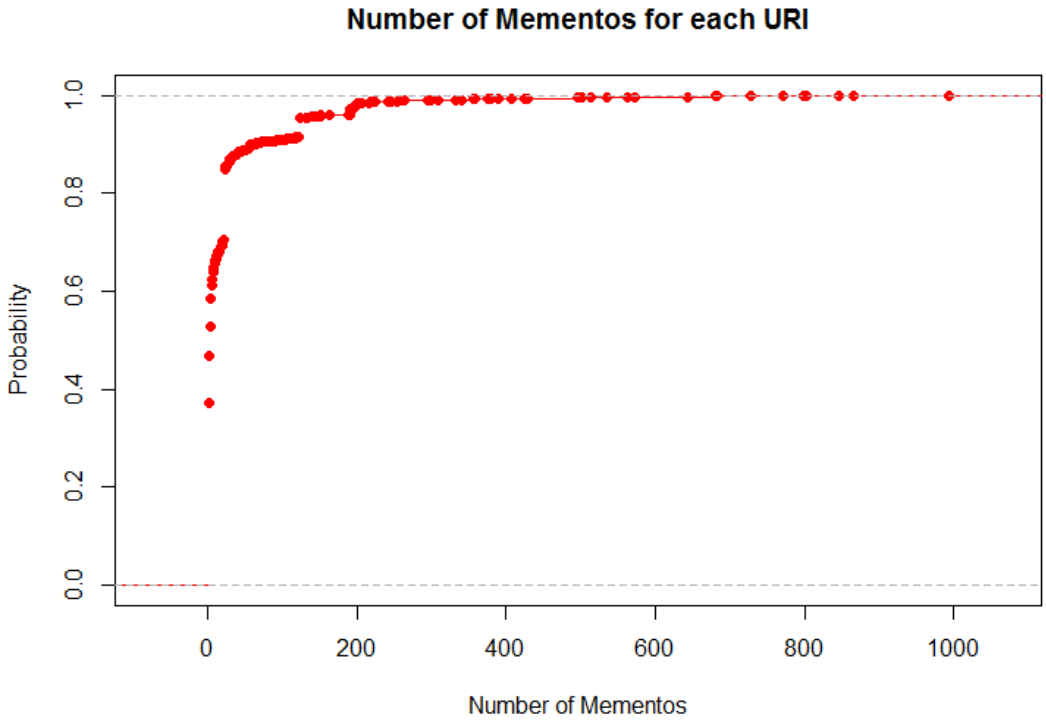

Fig. 2.1. CDF - Number of Mementos for each URI

2.3 Code Listing

```
1  import os
2  import datetime
3  import json
4
5  TIME_MAP_URL = "http://labs.mementoweb.org/timemap/json/"
6
7  f = open('status.txt', 'r+')
8  print datetime.datetime.now()
9  for line in f:
10         data = json.loads(line)
11         try:
12                 if data['index'] == 10001:
13                         print 'Program Executed'
14                         break
15                 finalURL=data['tweetURLData'][0]['finalURL']
16                 print '
                       ----------------------------------------------------------------------------------------------------------------------------------
                       ',
17                 os.system("wget --output-document=" + str(
                       data['index']) + " " + TIME_MAP_URL +
                       finalURL)
18         except:
19                 continue
20 print datetime.datetime.now()
```

Listing 2.1. Python program for fetching the TimeMaps.

```
1  import json
2  import os
3
4  w = open('mementoCount.txt', 'w+')
5  count = 1
6  while count < 10001:
7          path = 'labs/' + str(count)
8
9          if os.path.isfile(path):
10                 f = open(path, 'r+')
11                 if os.stat(path).st_size:
12                         try:
13                                 data = json.loads(f.read())
14                                 w.write(str(len(data['
                                       mementos']['list'])))
15                                 w.write('\n')
16                                 f.close()
17                         except KeyError:
18                                 w.write(str(len(data['
                                       timemap_index'])))
19                                 w.write('\n')
20                                 f.close()
21         count += 1
22         print count
23 w.close()
```

Listing 2.2. Python program for finding the number of mementos in each of the TimeMaps.

```
1 dp1 <- read.table('c:/users/kahmed/desktop/mementoCount.txt'
      , header=FALSE)
2 datapoint1 <- dp1[,1]
3 X1 = rnorm(sort(datapoint1))
4 P1 = ecdf(datapoint1)
5 plot(P1, col="red", xlab="Number of Mementos", ylab="
      Probability", main = "Number of Mementos for each URI")
```

Listing 2.3. R program for generating the CDF for number of mementos

3

Question 3

3.1 Question

- Using 20 links that have TimeMaps
  - With ¿= 20 mementos
  - Have existed >= 2 years (i.e., Memento-Datetime of first memento is April XX, 2013 or older)
  - Note: select from Q1/Q2 links, else choose them by hand
- For each link, create a graph that shows Jaccard Distance, relative to the first memento, through time
  - x-axis: continuous time, y-axis: Jaccard Distance relative to the first memento

3.2 Solution

- I manually searched for the 20 URIs that satisfied the condition.
- Then, I ran each of the URIs in the 20 TimeMaps through boilerpipe to retrieve the text. This would help me fetch the unigrams that I would then use to calculate the Jaccard Distance with respect to the first memento.
- I modified the script from question one to calculate the Jaccard Distance.
- Below are the graphs illustrating the Jaccard Distance relative to the first memento.

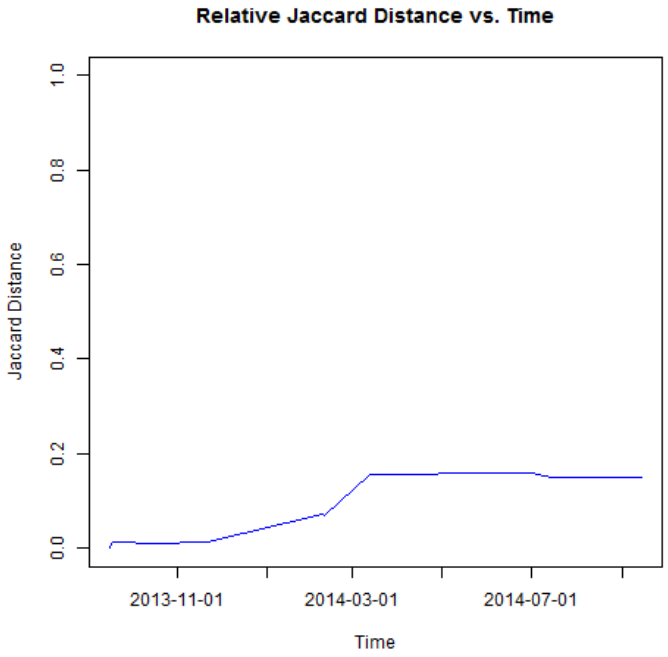

Fig. 3.1. Jaccard Distance relative to first memento for URI# 1

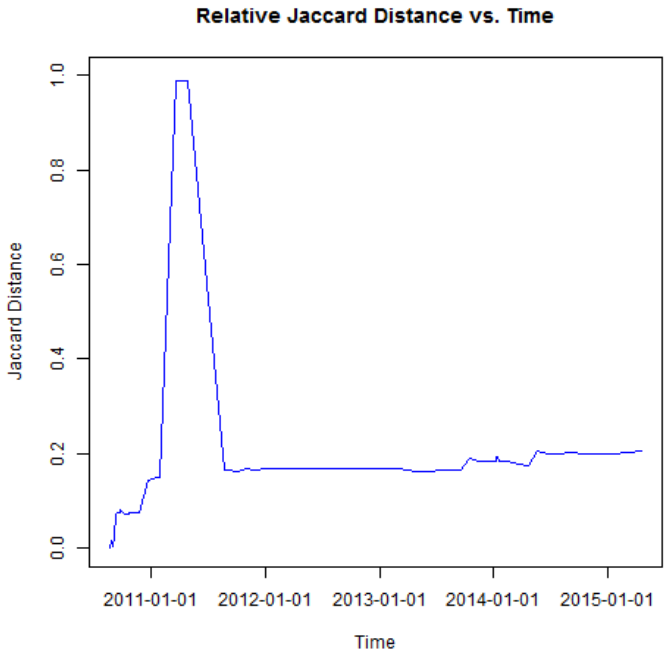

Fig. 3.2. Jaccard Distance relative to first memento for URI# 2

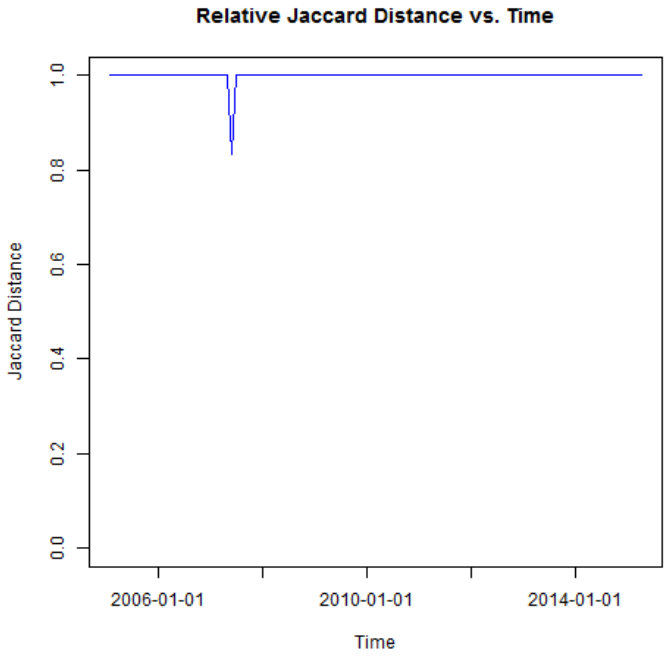

Fig. 3.3. Jaccard Distance relative to first memento for URI# 3

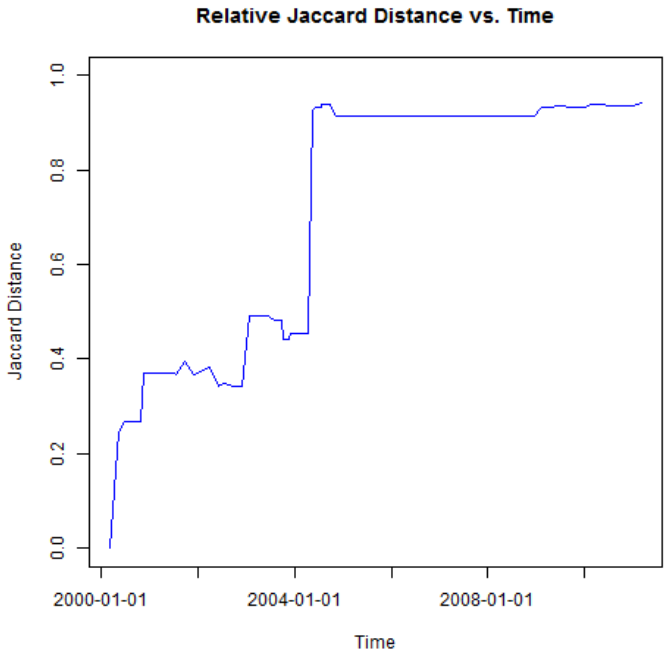

Fig. 3.4. Jaccard Distance relative to first memento for URI# 4

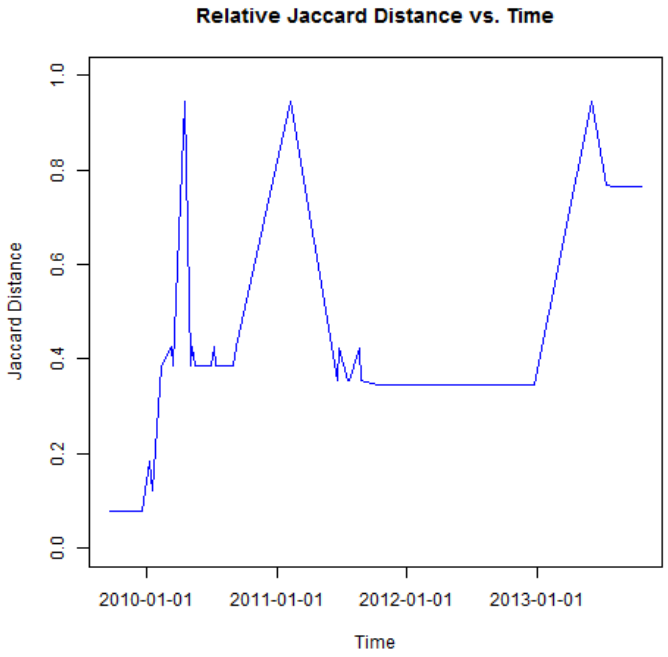

Fig. 3.5. Jaccard Distance relative to first memento for URI# 5

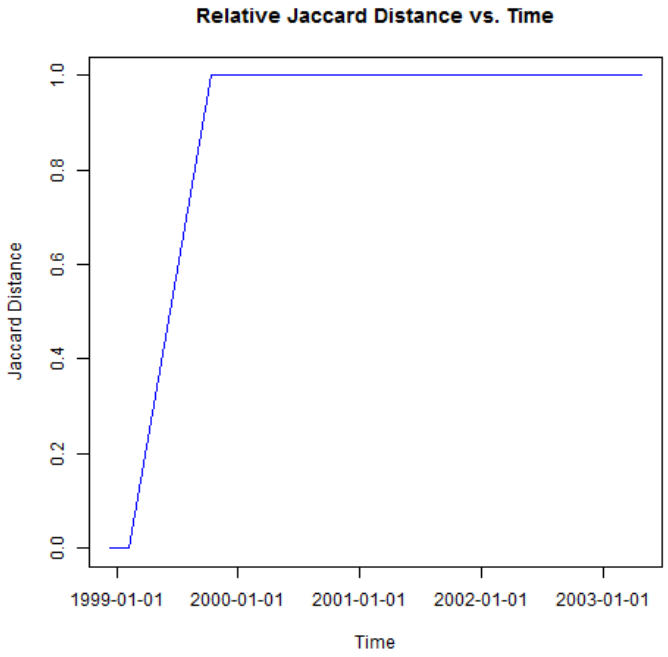

Fig. 3.6. Jaccard Distance relative to first memento for URI# 6

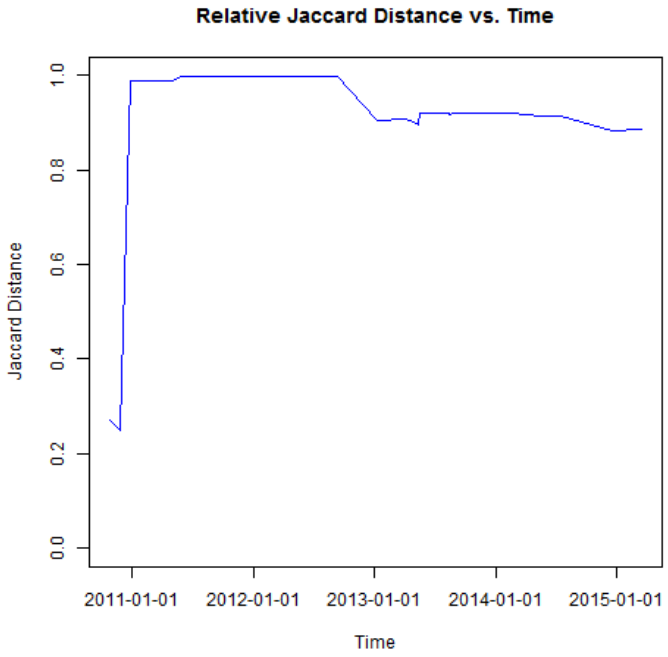

Fig. 3.7. Jaccard Distance relative to first memento for URI# 7

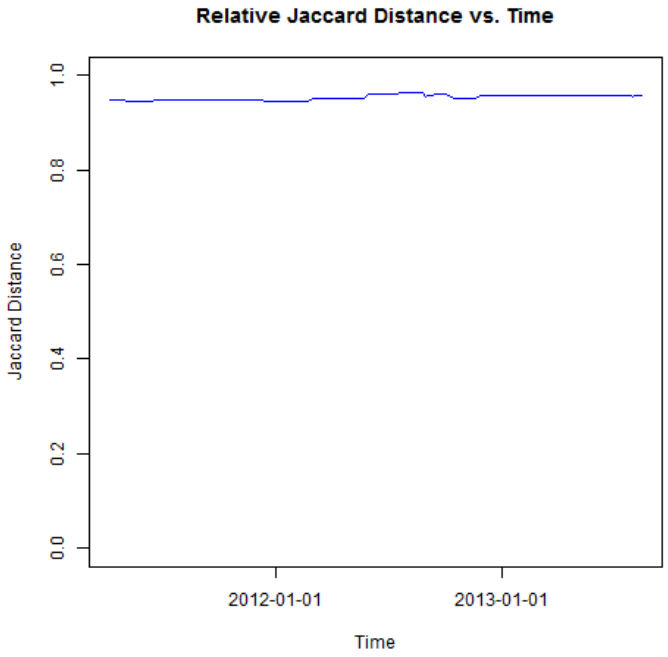

Fig. 3.8. Jaccard Distance relative to first memento for URI# 8

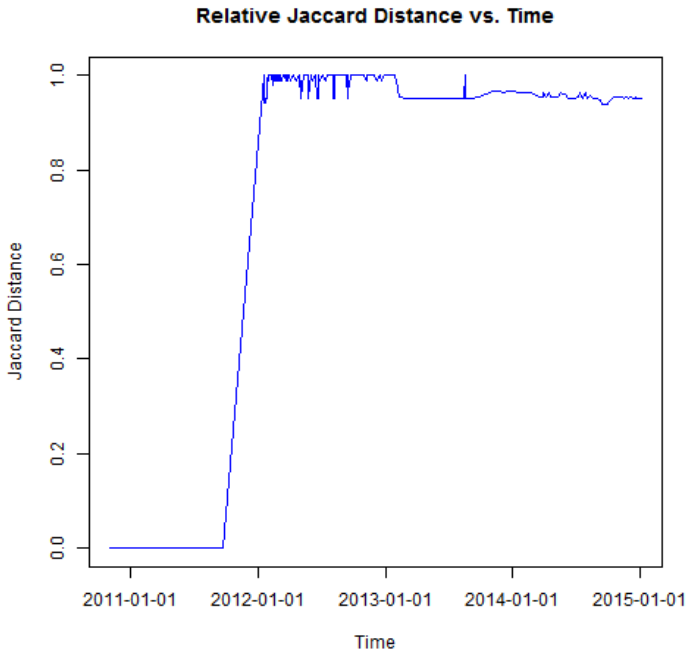

Fig. 3.9. Jaccard Distance relative to first memento for URI# 9

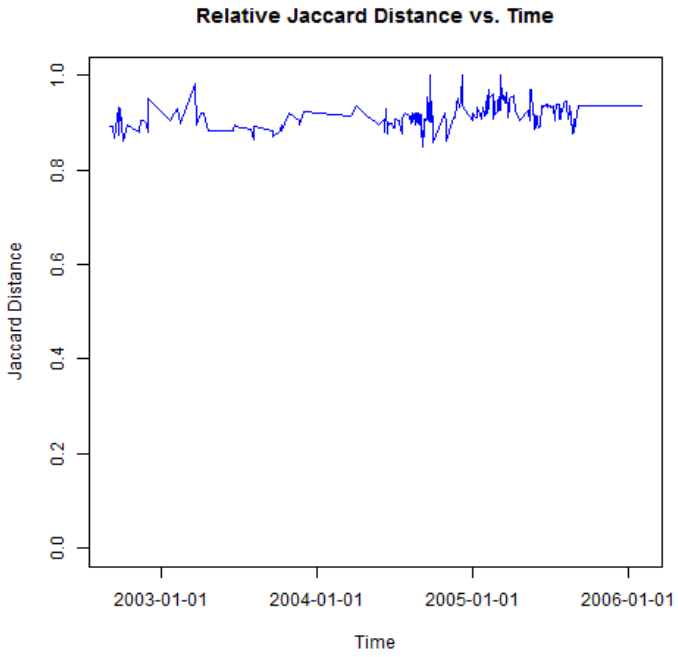

Fig. 3.10. Jaccard Distance relative to first memento for URI# 10

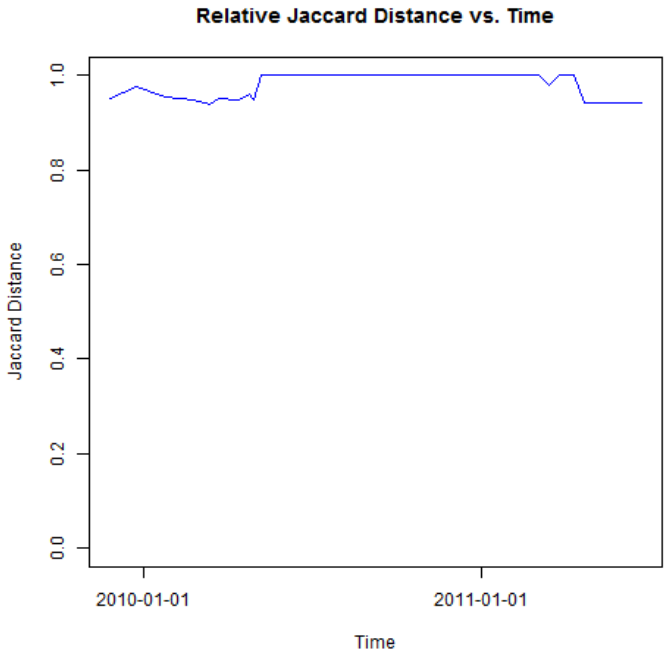

Fig. 3.11. Jaccard Distance relative to first memento for URI# 11

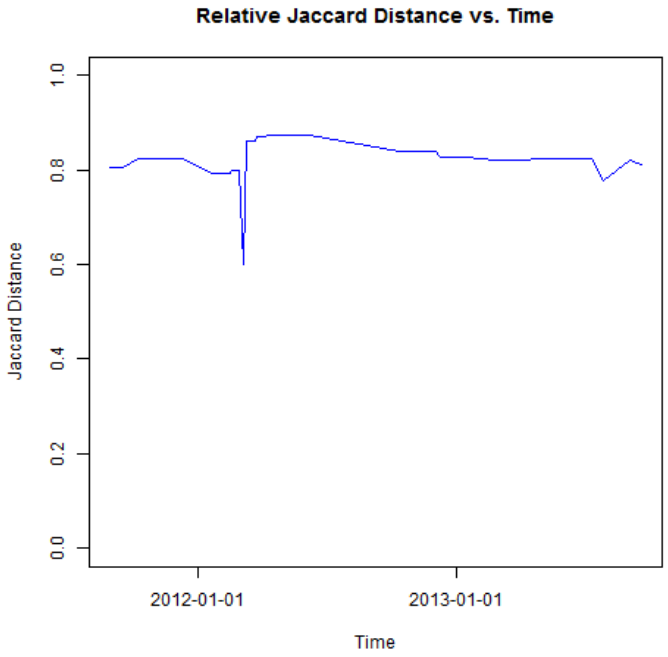

Fig. 3.12. Jaccard Distance relative to first memento for URI# 12

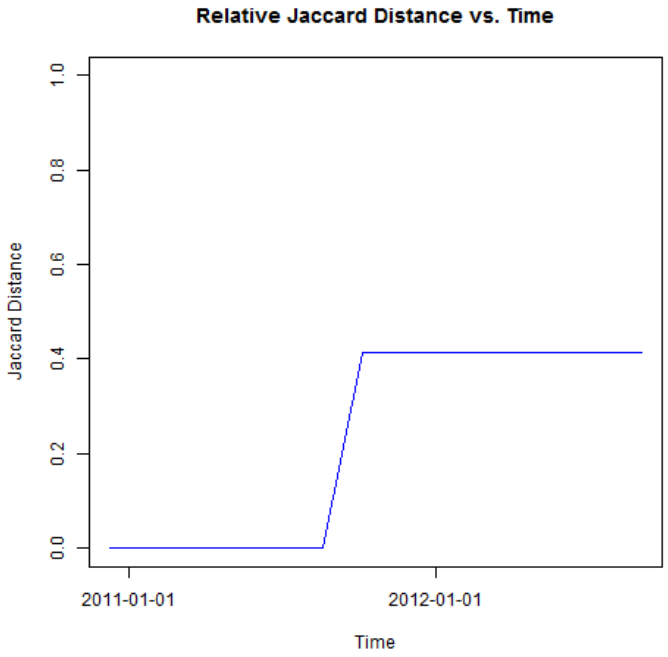

Fig. 3.13. Jaccard Distance relative to first memento for URI# 13

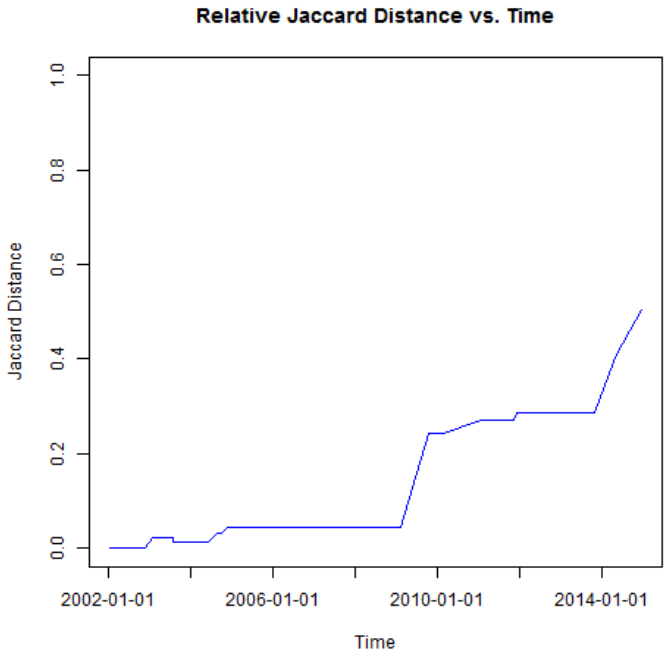

Fig. 3.14. Jaccard Distance relative to first memento for URI# 14

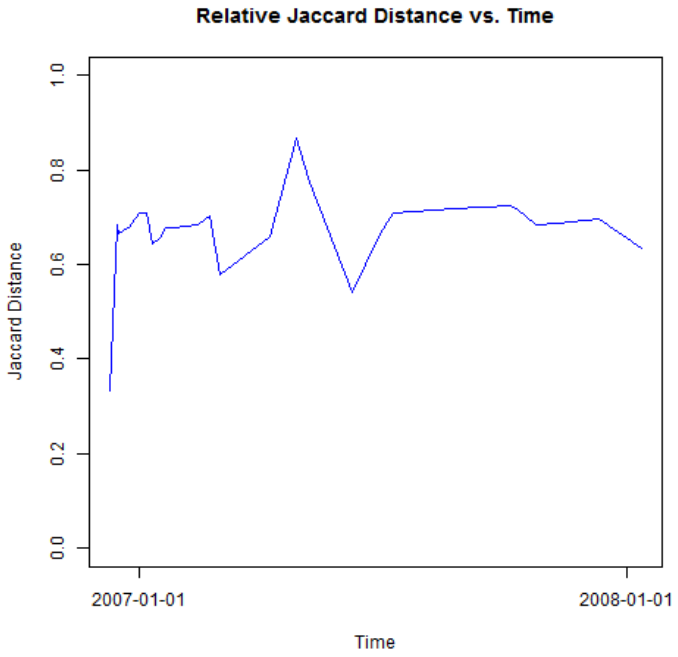

Fig. 3.15. Jaccard Distance relative to first memento for URI# 15

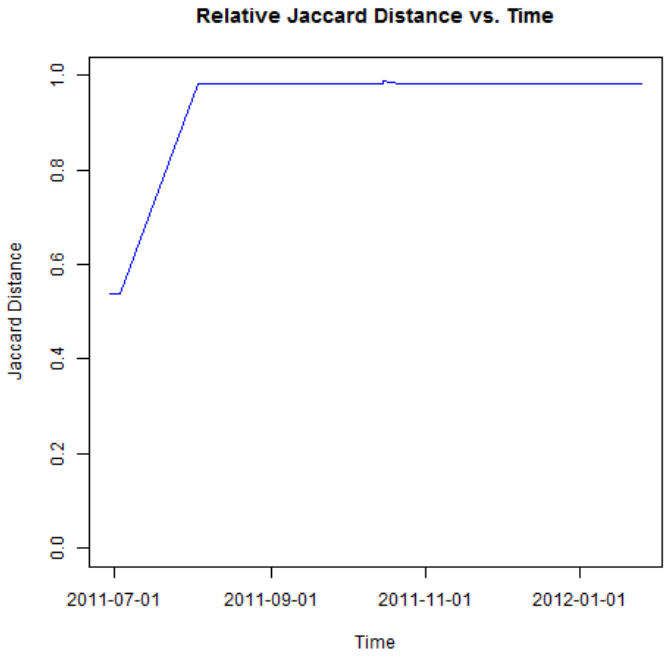

Fig. 3.16. Jaccard Distance relative to first memento for URI# 16

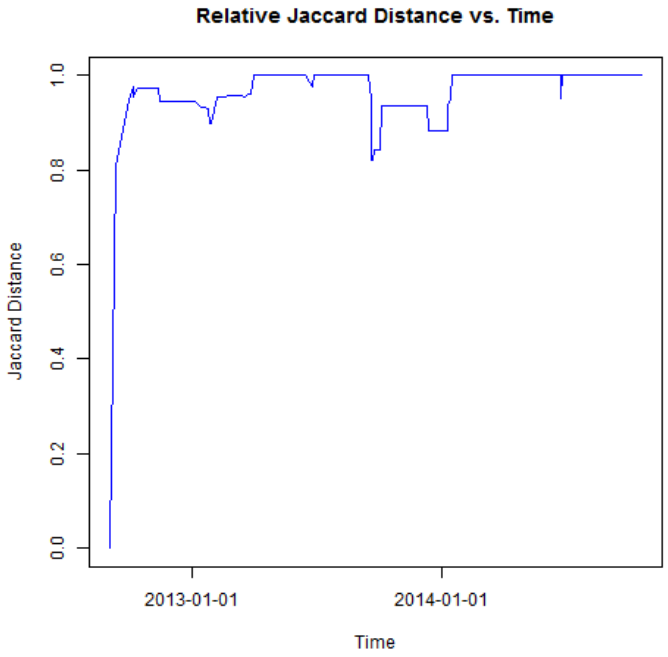

Fig. 3.17. Jaccard Distance relative to first memento for URI# 17

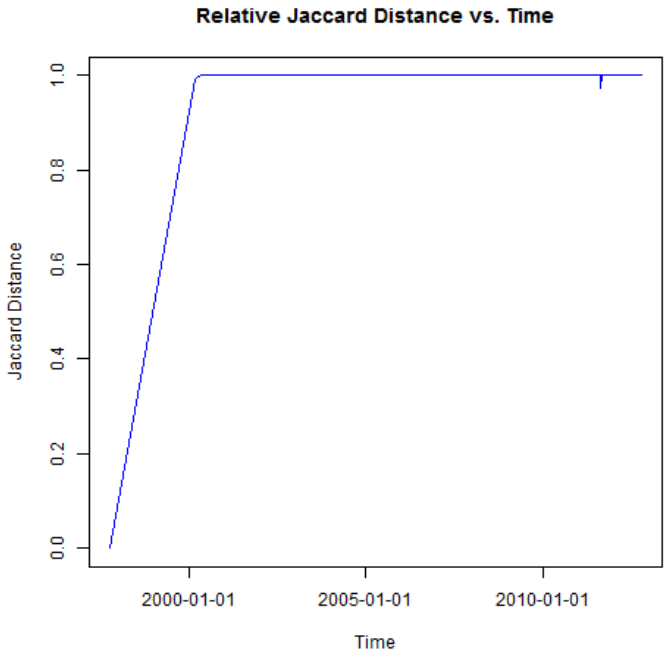

Fig. 3.18. Jaccard Distance relative to first memento for URI# 18

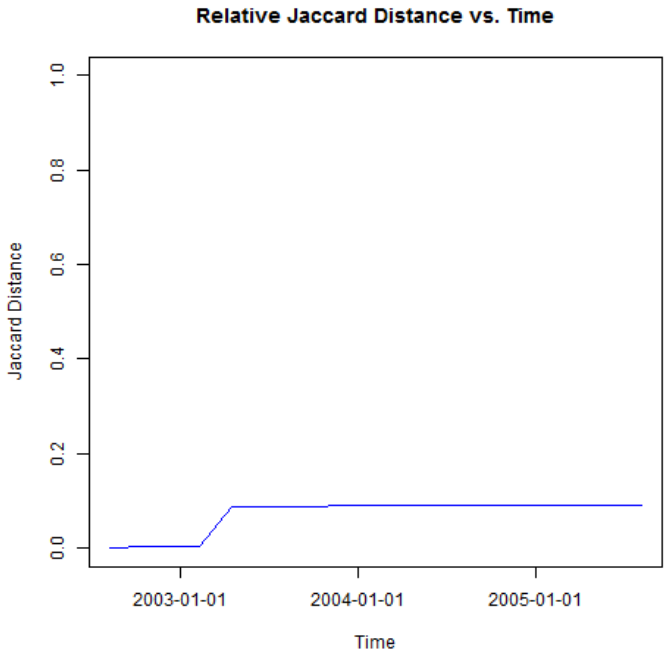

Fig. 3.19. Jaccard Distance relative to first memento for URI# 19

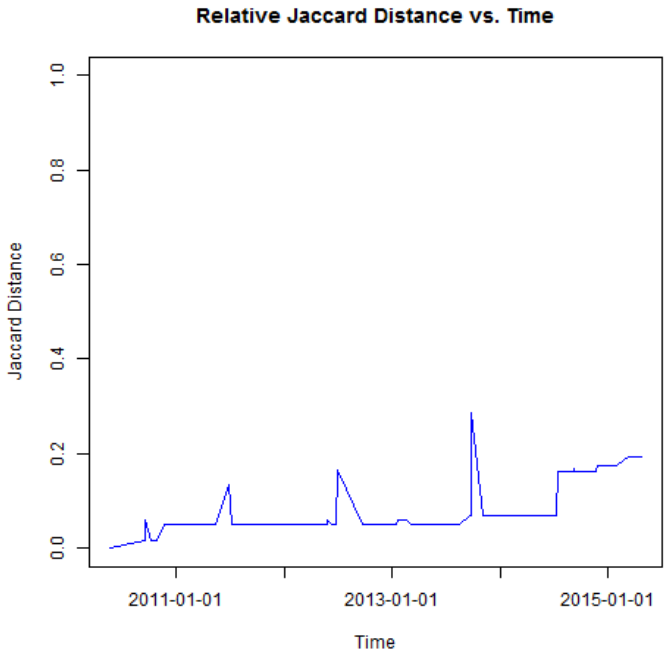

Fig. 3.20. Jaccard Distance relative to first memento for URI# 20

3.3 Code Listing

```
import json
import os
from boilerpipe.extract import Extractor

count = 1
while count < 10001:
        if count == 182 or count == 714 or count == 1106 or
            count == 1200 or count == 1417 or count == 2077
            or count == 2168 or count == 2235 or count ==
            2338 or count == 2604 or count == 5209 or count
            == 5986 or count == 6591 or count == 6969 or
            count == 7861 or count == 8145 or count == 8827
            or count == 9093 or count == 9548 or count ==
            9613:
                print count
                path = 'labs/' + str(count)

                if os.path.isfile(path):
                        f = open(path, 'r+')
                        if os.stat(path).st_size:
                                data = json.loads(f.read())
                                if not os.path.exists(str(
                                    count)):
                                        os.makedirs(str(
                                            count))

                                filecounter = 1
                                link = str(count) + '
                                    _datetime.txt'
                                g = open(link, 'w')
                                for url in data['mementos'][
                                    'list']:
                                        extractor =
                                            Extractor(
                                            extractor='
                                            DefaultExtractor
                                            ', url=url['uri'
                                            ])
                                        extracted_text =
                                            extractor.
                                            getText()
                                        g.write(str(url['
                                            datetime']))
                                        g.write('\n')
                                g.close()
                                filecounter += 1
```

```
28                              f.close()
29          count += 1
```

Listing 3.1. Python program for fetching boilerpipe content for each memento in the TimeMap.

```
1  from nltk.util import ngrams
2  import os
3
4  MAX_NGRAMS=1
5
6  def jaccardDistance(list1, list2):
7          union = set(list1).union(list2)
8          intersect = set(list1).intersection(list2)
9          if len(union) == 0:
10                 return 0
11         dist = (len(union) - len(intersect)) * 1.0 / len(
               union)
12         return dist
13
14 def getNGrams(fileName):
15         f = open(fileName, 'r+')
16         unigramList=[]
17         notEntered = True
18         #print fileName
19         for line in f:
20                 notEntered = False
21                 unigram = ngrams(line.split(), MAX_NGRAMS)
22                 for grams in unigram:
23                         unigramList.append(grams)
24                 #print len(unigramList)
25                 #print unigramList
26                 if len(unigramList) == 0:
27                         print fileName
28                         return unigramList,True
29         f.close()
30         if notEntered:
31                 return unigramList, True
32         return unigramList, False
33
34 def fileExists(fileName):
35         return os.path.isfile(fileName)
36 count = 1
37
38 while count < 10001:
39         if os.path.exists(str(count)):
40                 print str(count)
41                 fileCounter = 1
42                 fileName1 = str(count) + '/' + str(
                       fileCounter) + '.txt'
43                 if fileExists(fileName1):
44                         unigramListBaseline=[]
45                         shouldExit = False
```

```
46                          unigramListBaseline, shouldExit =
                                getNGrams(fileName1)
47                          #print shouldExit
48                          if shouldExit:
49                                  fileCounter += 1
50                                  fileName1 = str(count) + '/'
                                        + str(fileCounter) + '.
                                        txt'
51                                  unigramListBaseline=[]
52                                  shouldExit = False
53                                  unigramListBaseline,
                                        shouldExit = getNGrams(
                                        fileName1)
54                                  print 'exit1'
55                          #print unigramListBaseline
56                          uni = open(str(count) + '_jaccard.
                                txt','w')
57                          while fileCounter < 1000:
58                                  fileCounter += 1
59                                  fileName2 = str(count) + '/'
                                        + str(fileCounter) + '.
                                        txt'
60                                  if fileExists(fileName2):
61                                          unigramListCurrent
                                                =[]
62                                          shouldExit = False
63                                          unigramListCurrent,
                                                shouldExit =
                                                getNGrams(
                                                fileName2)
64                                          #print shouldExit
65                                          if shouldExit:
66                                                  print 'exit2
                                                        '
67                                                  #count += 1
68                                                  uni.write(
                                                        str(1.0)
                                                        )
69                                                  uni.write('\
                                                        n')
70                                                  continue
71                                          #print str(len(
                                                unigramListBaseline
                                                )) + '\t' + str(
                                                len(
                                                unigramListCurrent
                                                ))
72
```

```
73                                                  uni.write(str(
                                                        jaccardDistance(
                                                        unigramListBaseline
                                                        ,
                                                        unigramListCurrent
                                                        )))
74                                                  uni.write('\n')
75
76                              uni.close()
77              count += 1
78
79  print 'Program Executed'
```

Listing 3.2. Python program for calculating the Jaccard Distance relative to first memento.

```
1 data <- read.table('C:/Users/kahmed/Desktop/q3/uri20.txt',
      sep="\t", colClasses=c("POSIXct", "numeric"))
2 png('C:/Users/kahmed/Desktop/q3/uri20.png')
3 p1 <- plot(data, type="l", col='blue', main="Relative
      Jaccard Distance vs. Time", xlab="Time", ylab="Jaccard
      Distance", xaxt="n", ylim=c(0,1))
4 axis.POSIXct(side=1, data$V1, format="%Y-%m-%d")
```

Listing 3.3. R program for generating plot for Jaccard Distance vs. Time

4

Question 4

4.1 Question

- Choose a news-related event
- Use twarc.py to collect 1000 tweets, every day for 5 different days
  - See: `https://github.com/edsu/twarc`
- For each day:
  - Create a wall
  - Build a tag/word cloud for each day
  - Create a map using GeoJSON & Github
    - `https://help.github.com/articles/mapping-geojson-files-on-github/`
- Discuss in detail lessons learned, experiences, etc.

4.2 Solution

- I chose 'Google Fi' as the topic for this question.
- I installed the twarc package using 'sudo pip install twarc' on my ubuntu virtual machine.
- I created a script to fetch 1000 tweets and ran this script for five consecutive days.
- At the end of each day I followed the instructions of the author of the library to create the wall, word cloud, GeoJSON.
- As I progressed to fetch the tweets for the fourth day I noticed it took longer to fetch tweets. I suspected this could be attributed to the reduction in the discussion about this topic.
- Some of the tools provided by twarc are powerful.
- The wall displays the tweets as a html which I didn't think was of much help as there are multitude of websites that provide this facility of displaying tweets based on search parameters.

- I felt that the word cloud utility was a powerful feature that illustrates the words used in the tweet and the size of the word changes based on the frequency of usage in each of the tweets.
- The geojson utility would be a good tool to visualize the contributors and their location if the geo-location is shared along with the tweet. But due to the limited tweets that had the co-ordinates it would be too premature to make more sense of the feature.
- The html page generated by the wall utility is a good example of a file which would be boilerpipe successful. It has body content which can be extracted.
- The html page generated by the wordcloud utility is a perfect example of a file which would not be boilerpipe successful because it has no body content but receives all its data through scripts.
- I committed the files on github and the GeoJSON can be viewed using github pages from the following URL `https://github.com/maxbizarre/cs851-s15/blob/master/assignment4/twarc/tweetDay1.geojson` and `https://github.com/maxbizarre/cs851-s15/blob/master/assignment4/twarc/tweetDay2.geojson`.

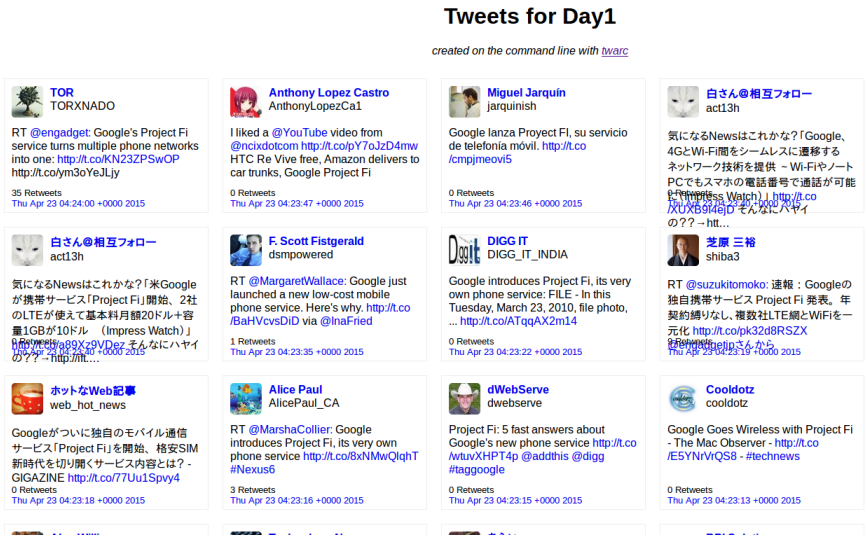

Fig. 4.1. Wall - Day 1

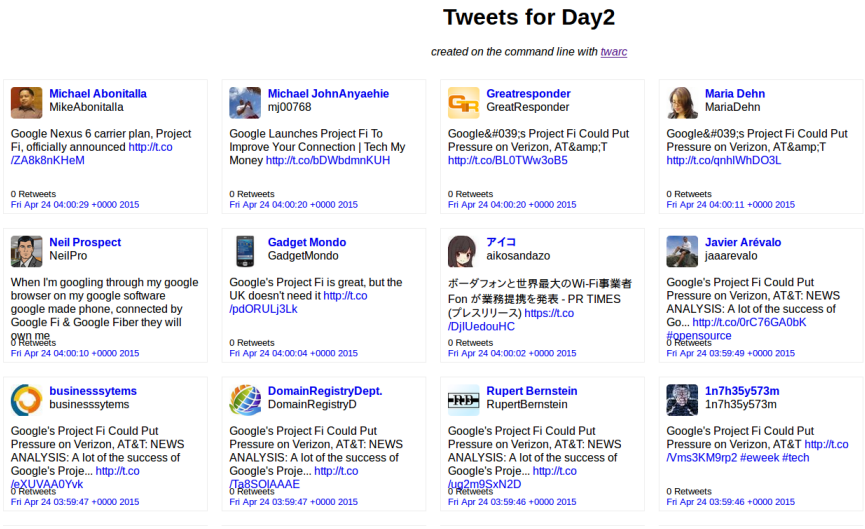

Fig. 4.2. Wall - Day 2

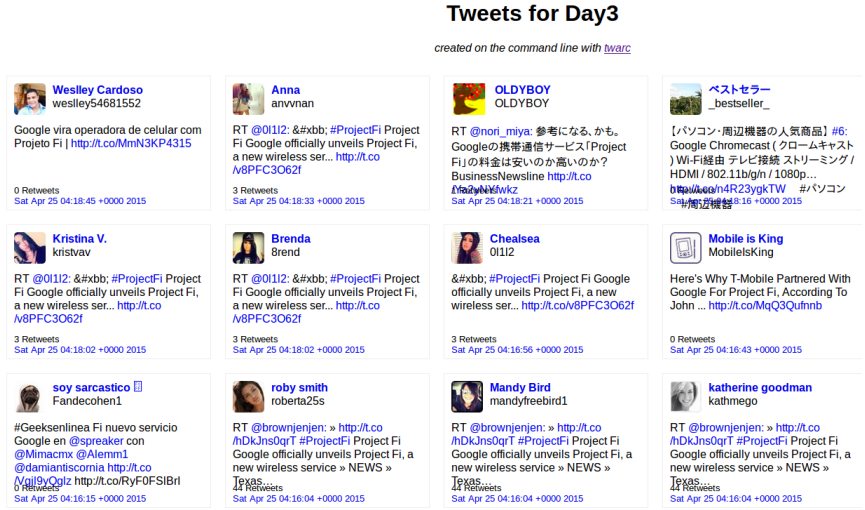

Fig. 4.3. Wall - Day 3

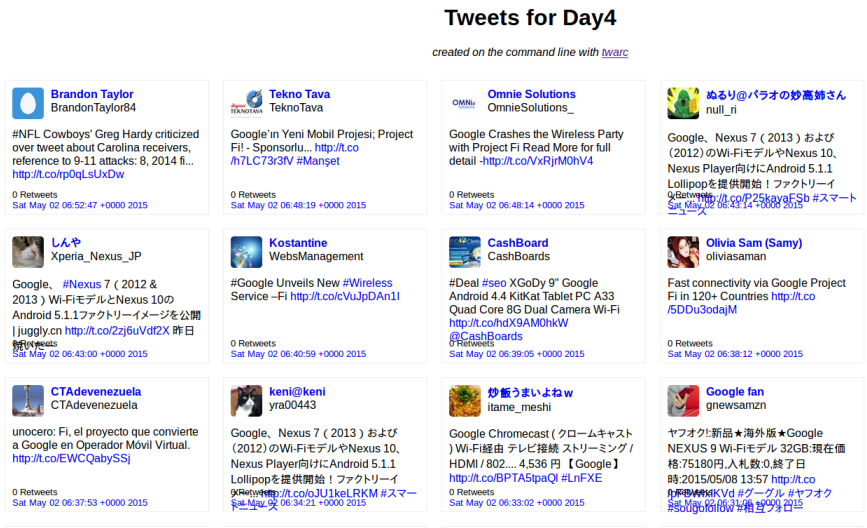

Fig. 4.4. Wall - Day 4

Fig. 4.5. Wall - Day 5

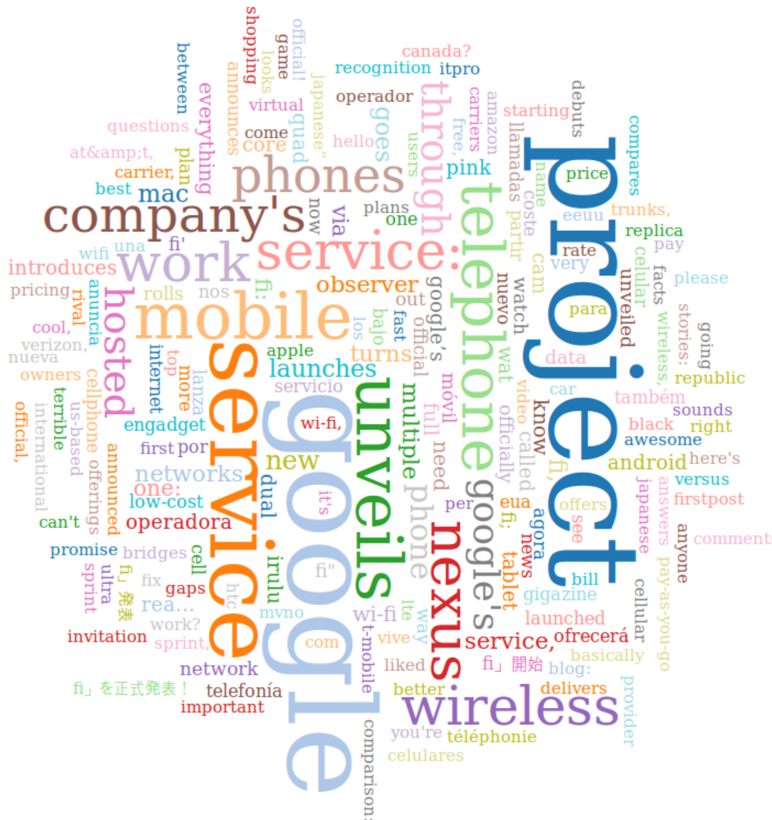

Fig. 4.6. Word Cloud - Day 1

Fig. 4.7. Word Cloud - Day 2

Fig. 4.8. Word Cloud - Day 3

Fig. 4.9. Word Cloud - Day 4

Fig. 4.10. Word Cloud - Day 5

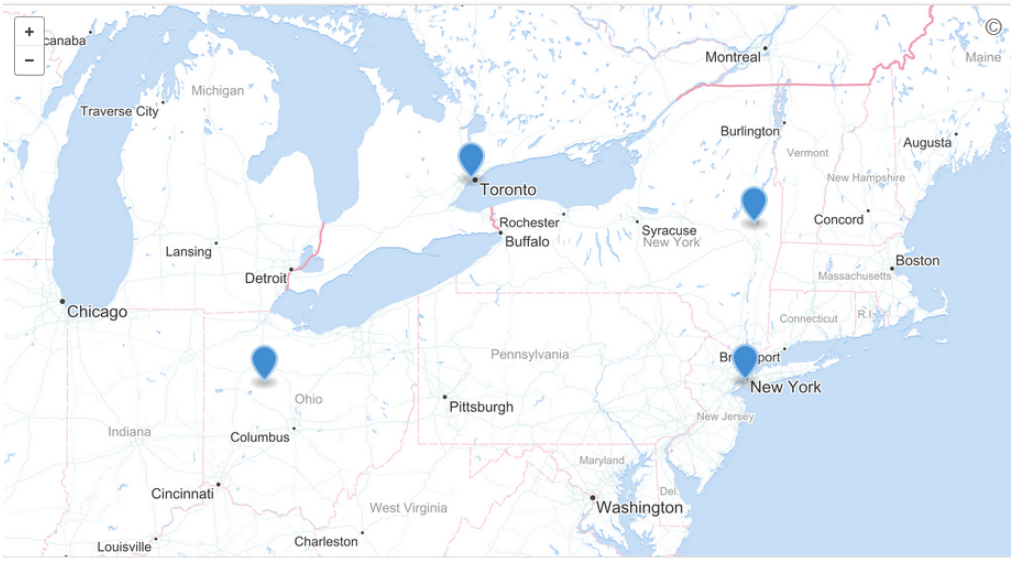

Fig. 4.11. Geographical Location - Day 1

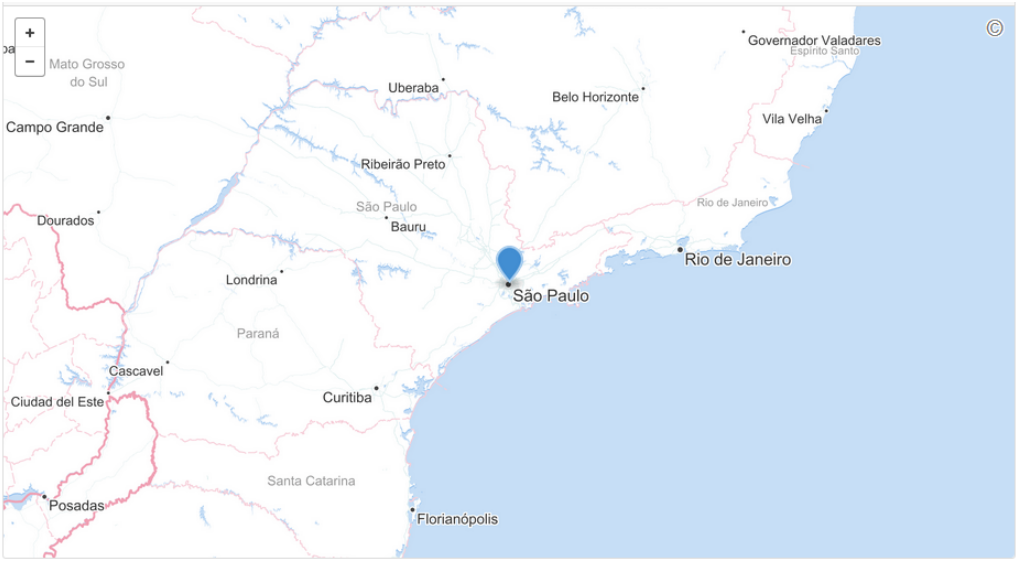

Fig. 4.12. Geographical Location - Day 2

4.3 Code Listing

```
1  from twarc import Twarc
2  import json
3
4  CONSUMER_KEY = 'CfHUyBhlMaLv5Mn8r2IziXpLs'
5  CONSUMER_SECRET = '
       PqqtbhbyNb5mcJ2dHkSIT2wupOMuEqfSINGYvV8KDIOPuqgDkN'
6  ACCESS_TOKEN = '29202483-
       qK6twPLeurVc8Ls8zBxdFtaFGyzm76LUBbtXOMMk1'
7  ACCESS_TOKEN_SECRET = '
       aOIFdI1TVJjsIPWNO1rAFx2IECzVSCPY4kOnEKBA0pCdA'
8
9  w = open('tweetDay5.json', 'w')
10
11 t = Twarc(CONSUMER_KEY, CONSUMER_SECRET, ACCESS_TOKEN,
       ACCESS_TOKEN_SECRET)
12 count = 1
13 for tweet in t.search("google fi"):
14         w.write(json.dumps(tweet))
15         w.write('\n')
16         print count
17         count += 1
18         if count >1000:
19                 break
```

Listing 4.1. Python program for fetching tweets using twarc.

References

1. Cdf in r. http://stats.stackexchange.com/questions/30858/how-to-calculate-cumulative-distribution-in-r.
2. n-grams in python. http://stackoverflow.com/questions/13423919/computing-n-grams-using-python.
3. twarc. https://github.com/edsu/twarc.